



Big Data Search - The ELK Trifecta

Demo: Searching Your Spend

Melanie Akwule

Agenda



**Elastic Overview:
Elastic and its
open source tools**


**Demo Summary:
Searching your spend**

**Getting Started:
The ELK trifecta**



Elastic Overview

elastic and its open source tools



Elastic Overview

The ELK Trifecta



kibana

Visualization platform




elasticsearch

Search and analytics engine



logstash

Data collection, enrichment,
and transportation pipeline



Demo Summary

searching your spend

Today's Demo

- Builds out a dashboard that highlights insightful metrics
- Walks you through how to build a visualization
- Creates an elasticsearch index called "womsa" to analyze each record
- Maintains out of the box configuration
- Ingests a ~28K record file of mock financial data
- Key Parameters: Full Date Timestamp, Transaction, Description, Account Name, & Category



kibana
Visualization platform



elasticsearch
Search and analytics engine



logstash
Data collection, enrichment,
and transportation pipeline



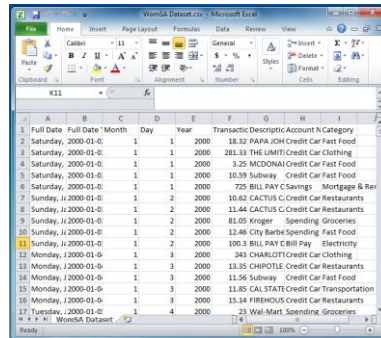
Getting Started

trying it at home

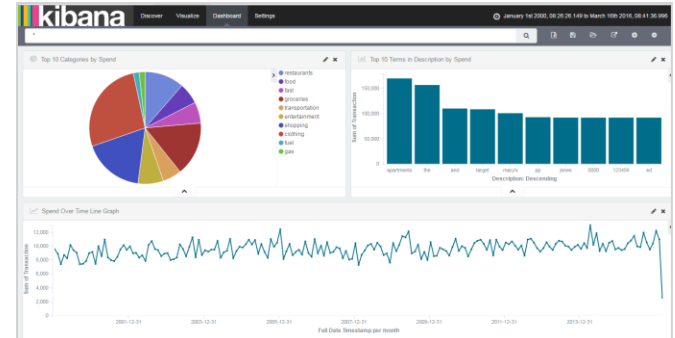
What you'll find on GitHub...




This Presentation



Dataset used to build the Dashboard




The dashboard from Kibana



Logstash configuration file

Getting Started *elasticsearch*



- 1 • Java [Install](#)
• Setting JAVA_HOME Variable [Documentation](#)


```
JAVA_HOME environment variable must be set
Press any key to continue . . .
```

- 2 • Elasticsearch 2.2.0 [Download](#)
• Getting Started [Documentation](#)

- 3 • Start ES: \elasticsearch-2.2.0\bin>elasticsearch
• Helpful CMD [Documentation](#)

```
C:\Users\...>\Desktop\elasticsearch-2.2.0\bin>elasticsearch
[2016-03-10 11:18:21.114][WARN ][bootstrap] jvm uses the client
vm, make sure to run 'java' with the server vm for best performance by adding '-
server' to the command line
[2016-03-10 11:18:21.395][INFO ][node]
on[2.2.0], pid[7212], build[8ff36d1/2016-01-27T13:32:39Z]
[2016-03-10 11:18:21.395][INFO ][node]
alizing ...
[2016-03-10 11:18:22.218][INFO ][plugins]
es [lang-expression, lang-groovy], plugins [], sites []
[2016-03-10 11:18:22.265][INFO ][node]
[1] data paths, mounts [[Windows? (C:)]], net usable_space [120gb], net total_s
pace [238.4gb], spins? [unknown], types [NTFS]
[2016-03-10 11:18:22.265][INFO ][node]
size [989.8mb], compressed ordinary object pointers [unknown]
[2016-03-10 11:18:24.335][INFO ][node]
alized
[2016-03-10 11:18:24.335][INFO ][node]
ing ...
[2016-03-10 11:18:24.740][INFO ][transport]
ch_address [127.0.0.1:9300], bound_addresses [127.0.0.1:9300]
[2016-03-10 11:18:24.756][INFO ][discovery]
icsearch-y8KHoMypI-OVSxUuSkXUHA
[2016-03-10 11:18:28.840][INFO ][cluster.service]
aster <Grasshopper> <y8KHoMypI-OVSxUuSkXUHA> [127.0.0.1] <127.0.0.1:9300>, reason:
sen-disco-join[electd_as_master, {0} joins received]
[2016-03-10 11:18:28.949][INFO ][gateway]
read [0] indices into cluster_state
[2016-03-10 11:18:29.152][INFO ][http]
ch_address [127.0.0.1:9200], bound_addresses [127.0.0.1:9200]
[2016-03-10 11:18:29.168][INFO ][node]
ed
```

- 4 • Sense Chrome Plugin [Download](#)
• Query DSL [Documentation](#)




- 5 • Run ES: <http://localhost:9200/>
• List of ES Indices: http://localhost:9200/_cat/indices?v
• Looking inside an index: http://localhost:9200/test2/_search?pretty

```
{
  "name": "Grasshopper ",
  "cluster_name": "elasticsearch",
  "version": {
    "number": "2.2.0",
    "build_hash": "8ff36d139e16f8720f2947ef62c8167a888992fe",
    "build_timestamp": "2016-01-27T13:32:39Z",
    "build_snapshot": false,
    "lucene_version": "5.4.1"
  },
  "tagline": "You Know, for Search"
}
```

| health | status | index | pri | rep | docs.count | docs.deleted | store.size | pri.store.size |
|--------|--------|---------|-----|-----|------------|--------------|------------|----------------|
| yellow | open | movies | 5 | 1 | 1 | 0 | 4.4kb | 4.4kb |
| yellow | open | test2 | 5 | 1 | 1231 | 0 | 626.7kb | 626.7kb |
| yellow | open | .kibana | 1 | 1 | 4 | 1 | 19.8kb | 19.8kb |
| yellow | open | test1 | 5 | 1 | 6155 | 0 | 2.6mb | 2.6mb |

Getting Started *logstash*



- 1 • JRuby [Install](#)

```
D:\logstash\bin>logstash -e 'input { stdin { } } out;
Unable to find JRuby.
If you are a user, this is a bug.
If you are a developer, please run 'rake bootstrap'.
```

- 2 • Logstash 2.2.2 [Download](#)
 • Getting Started [Documentation](#)

- 3 • Enter Standard I/O Command: logstash -e 'input { stdin { } } output { stdout { } }'
 • Should be able to type and have it echoed back

```
C:\Users\... \Desktop\logstash-2.2.2\logstash-2.2.2\bin>logstash -e 'input
{ stdin { } } output { stdout { } }'
io/console not supported; tty will not be manipulated
Settings: Default pipeline workers: 4
Logstash startup completed
```

```
C:\Users\... \Desktop\logstash-2.2.2\logstash-2.2.2\bin>logstash -e 'input
{ stdin { } } output { stdout { } }'
io/console not supported; tty will not be manipulated
Settings: Default pipeline workers: 4
Logstash startup completed
i typed this
2016-03-13T17:03:40.369Z i typed this
and then now i typed this as well
2016-03-13T17:03:51.273Z and then now i typed this as well
```

- 4 • Preparing a Configuration File [Documentation](#)
 • Configuring for CSV [Documentation](#)
 • List of Input [Plugins](#)
 • List of Output [Plugins](#)


```
1 input {
2   file {
3     path => "C:/Users/... /Desktop/WomSA Search Demo/Dashboard/*.csv"
4     type => "womsa"
5     start_position => "beginning"
6   }
7 }
8
9 filter {
10  csv {
11    columns => ["Full Date", "Full Date Timestamp", "Month", "Day", "Year",
12    separator => ","
13  }
14
15  mutate {
16    convert => [ "Transaction", "float" ]
17  }
18 }
19
20 output {
21  elasticsearch {
22    action => "index"
23    hosts => "127.0.0.1"
24    index => "womsa"
25    workers => 1
26  }
27  stdout {
28    codec => json
29  }
30 }
```

- 5 • Import File: logstash -f <filename>.conf
 • Troubleshooting Field Format [Documentation](#)
 • Troubleshooting Data Ingestion [Documentation](#)

```
\bin>logstash -f womsa.conf
```


Getting Started

kibana



1

2

3

4

Start ES service

`\bin>elasticsearch`

Start kibana service

`\bin>kibana`

Configure Kibana
<http://localhost:5601/>

1

- Kibana 4.4.2 [Download](#)
- Getting Started [Documentation](#)

2

- Start ES: `\elasticsearch-2.2.0\bin>elasticsearch`

3

- Start Kibana: `\kibana-4.4.2\bin>kibana`

```
Microsoft Windows [Version 6.1.7601]
Copyright (c) 2009 Microsoft Corporation. All rights reserved.

C:\Users\> \Desktop\kibana-4.4.2-windows\bin>kibana
'WHERE' is not recognized as an internal or external command,
operable program or batch file.
log [16:29:21.124] [info][status][plugin:kibana] Status changed from uninitia
lized to green - Ready
log [16:29:21.155] [info][status][plugin:elasticsearch] Status changed from
uninitialized to yellow - Waiting for Elasticsearch
log [16:29:21.170] [info][status][plugin:kbn_vislib_vis_types] Status change
d from uninitialized to green - Ready
log [16:29:21.170] [info][status][plugin:markdown_vis] Status changed from u
ninitialized to green - Ready
log [16:29:21.248] [info][status][plugin:metric_vis] Status changed from uni
nitialized to green - Ready
log [16:29:21.264] [info][status][plugin:spyModes] Status changed from unin
itialized to green - Ready
log [16:29:21.280] [info][status][plugin:statusPage] Status changed from uni
nitialized to green - Ready
log [16:29:21.326] [info][status][plugin:table_vis] Status changed from unin
itialized to green - Ready
log [16:29:21.342] [info][listening] Server running at http://0.0.0.0:5601
log [16:29:21.389] [info][status][plugin:elasticsearch] Status changed from
yellow to green - Kibana index ready
```

4

- Change the name of your index to match ES
- Map time field – "yyyy-mm-ddThh:mm:ss"

☒ Index contains time-based events

☐ Use event times to create index names [DEPRECATED]

Index name or pattern

Patterns allow you to define dynamic index names using * as a wildcard. Example: logstash-*

logstash-*

☐ Do not expand index pattern when searching (Not recommended)

By default, searches against any time-based index pattern that contains a wildcard will automat range.

Searching against the index pattern `logstash-*` will actually query elasticsearch for the specific r

Unable to fetch mapping. Do you have indices matching the pattern?

☒ Index contains time-based events

☐ Use event times to create index names [DEPRECATED]

Index name or pattern

Patterns allow you to define dynamic index names using * as a wildcard. Example: logstash-*

womsa

Time-field name refresh fields

Full Date Timestamp

Create

Let's check it out..